

BRIEFING · SEPTEMBER 2026

Artificial Intelligence in *UK Central Government Services*

What the evidence actually says, what the oversight bodies expect, and what can safely be done now. The Department for Work and Pensions' Universal Credit decisions cover Great Britain; Northern Ireland runs its own system through the Department for Communities, and Scotland's devolved benefits sit with Social Security Scotland. Home Office asylum decisions are UK-wide.

Neil Catton

Independent Technology Strategist · Technology, Ethics & Human Purpose

NEILCATTON.COM

Prepared as a market briefing. Every figure is traced to a department's own published record, the National Audit Office, a select committee, a tribunal, official statistics or the statute book.

A note on the numbers

This briefing is built from what can be checked, and it names what was left out.

There is no single regulator for central government. Its use of artificial intelligence is overseen instead by the bodies that audit it, question it and hear appeals against it: the National Audit Office, the Public Accounts Committee, the Information Commissioner, and the tribunals. Every figure here traces to one of them, to official statistics, to the statute book, or to a department speaking on its own published record. No supplier material is used.

NAMED AND EXCLUDED: THE MOST QUOTABLE NUMBER IN THE SECTOR

Government's own estimate of £45bn a year from the "full potential digitisation" of public services is excluded. Its methodology note says it does not assess whether automation is feasible and sets no timeframe, and the full model was not released. The National Audit Office's May 2026 guide warns that "external evidence suggests that promises of greater productivity translating into substantial organisational financial savings are not yet borne out by real world experience". The £36bn automation component goes with it, as does a figure that 62% of administrative assistants' tasks are "routine": routine is not the same as automatable. Figures on DWP bank monitoring circulating on accountancy blogs are unsourced and excluded.

FLAGGED, NOT HIDDEN

An Upper Tribunal decision of 14 July 2026 on a Home Office asylum refusal is reported by The Guardian and others but was not on the tribunal's published database when it was searched for this briefing. It appears as an illustration, attributed to that reporting and stated as unproven. It does not carry the argument.

DATES AND COMPARATORS, BECAUSE THEY MOVE

DWP changed the comparator age group between its 2025 and 2026 assessments of its advances model, so no year-on-year claim about age is made. The National Audit Office's 2023 survey of AI adoption across government is primary but nearly three years old, and is not used.

Sources for each figure are given at the foot of the section in which it appears, and listed in full in section thirteen.

The wider market

Decisions at a scale no other sector makes, and a state publishing more about its algorithms than ever.

£9.5bn

Benefit overpaid through fraud and error in 2024-25, 3.3% of expenditure

1.3m

Universal Credit advances paid in 2025-26, worth £700m

152

Algorithmic transparency records on GOV.UK, 29 September 2026

Central government makes decisions about entitlement, status and money for millions of people a year. The National Audit Office puts benefit overpayments through fraud and error at 3.3 per cent of expenditure in 2024-25, £9.5bn, down from 3.6 per cent the year before. The Department for Work and Pensions paid 1.3 million Universal Credit advances in 2025-26, worth £700m, and estimates fraud and error on those advances at between £20m and £90m. The Public Accounts Committee says the public purse as a whole "is losing c.£55-£81bn a year to fraud and error".

The strength of the position is transparency. The Algorithmic Transparency Recording Standard is mandatory for central government departments and in-scope arm's length bodies. In January 2025 there were 33 records on the register. On 29 September 2026 there were 152 — close to a fivefold increase in twenty months. DWP has seven records and HMRC seven. The register is imperfect, and section five returns to what is missing from it, but no private sector in this series publishes anything comparable about the algorithms it runs.

WHY THIS MATTERS TO AN AI PROGRAMME

The question is no longer whether AI is in use. It is what the review around it can show.

Government's own account of its AI rests on the person who makes the final decision. That is the right place to put the weight. It is also the part of the system about which least is published.

NAO, *Tackling benefit overpayments due to fraud and error*, 22 October 2025 · DWP, *Effectiveness Assessment of Universal Credit Advances Model*, 9 July 2026 · PAC, 27 March 2026 · PAC, *Use of AI in Government*, 26 March 2025 · GOV.UK algorithmic transparency records, 29 September 2026.

The pressure point

The human review that demonstrably changes outcomes comes last, months after the decision.

In the three months from April to June 2026, the Social Security and Child Support tribunal cleared 20,000 cases at hearing. Of these, 60 per cent were overturned in favour of the claimant. For Personal Independence Payment the figure was 67 per cent; for Universal Credit, 45 per cent. In the immigration and asylum chamber, 38 per cent of the 7,500 asylum appeals determined in the quarter were allowed.

BENEFIT APPEALS, APRIL TO JUNE 2026

60% of decisions heard were overturned in the claimant's favour. Receipts rose 55% on a year to 46,000; the open caseload rose 63% to 129,000.

IMMIGRATION AND ASYLUM APPEALS

38% of asylum appeals determined were allowed. The chamber's open caseload rose 48% to 156,000.

The statistics record outcomes, not reasons, and a tribunal overturn is not proof that the first decision was wrong on the evidence the first decision-maker had. But the tribunal is the one human review in the system whose effect is measured and published every quarter. The caseworker's review, which is the review government points to when it defends its use of AI, has no published measure at all.

That is the pressure point for anyone who owns a department's AI programme. The law now asks whether human involvement in a decision is meaningful. The only review in the system that can show its effect is the one at the end — slower, more expensive, and reached only by the people who appeal.

Ministry of Justice, *Tribunal Statistics Quarterly: April to June 2026*, 10 September 2026.

What the sector is facing

Two departments under pressure to decide more, faster – and a model that shows what the review is up against.

The Department for Work and Pensions. The Public Accounts Committee reported in March 2026 that "the use of data analytics to tackle fraud and error is underdeveloped across government", and asked that all relevant uses of AI and machine learning be captured on the transparency register. DWP's Universal Credit advances model is the furthest along. DWP describes it as "the only fraud and error machine learning model currently deployed at scale in live service". It flags advance requests for extra checks, and it is, by DWP's measure, 2.5 times more effective than a random control group at identifying high-risk advances. The NAO estimates it saved £4.4m between May 2022 and its October 2025 report.

The Home Office. At the end of June 2026, 40,168 people were awaiting an initial asylum decision, 56 per cent fewer than a year earlier. Meanwhile the open asylum and protection caseload at the First-tier Tribunal rose from 6,319 in June 2022 to 87,450 in March 2026. The department piloted two AI tools for asylum decision-makers: Asylum Case Summarisation, which summarises interview transcripts, and Asylum Policy Search, which finds and summarises country policy information. In March 2026 a Home Office minister told Parliament: "APS has been rolled out and is accessible to all Asylum decision makers and ACS is awaiting operationalisation."

What DWP's own assessment shows

DWP publishes an annual assessment of the advances model, which is to its credit. The 2025-26 edition, published on 9 July 2026, reports how much more likely each group is to be referred by the model than a comparator group, and how much more likely a referral is to be correct.

Group	Relative likelihood of referral	Of a correct referral
Non-UK nationals	3.35 (2.27 in 2024-25)	1.04 (0.97)
Aged 45 to 54	1.71	0.89
Aged 55 to 65	1.36	0.58
Aged 66 and over	27.24 (0.1% of observations)	0.36
Couples	0.98	0.36

Non-UK nationals are referred at more than three times the rate of UK nationals, and those referrals are correct at the same rate. DWP's own conclusion is that "the model is not working as effectively as it could for this cohort". The referral ratio rose from 2.27 to 3.35 in a year. Older claimants are referred more often and are less likely to be correctly referred. DWP's assessment remains that "it remains reasonable and proportionate to continue operating the UC Advances model". The model has been retrained, "but further testing is required ahead of the retrained model going live".

PAC, 27 March 2026 · DWP, *Universal Credit Advances Model – Fairness Assessment*, 17 July 2025, and *Effectiveness Assessment*, 9 July 2026 (age comparator 25–34; nationality comparator UK nationals; couples comparator single claimants) · NAO, 22 October 2025 · Home Office, *Immigration system statistics, year ending June 2026*, 27 August 2026 · Home Office, written answer UIN 121674, 24 March 2026.

The challenges of using AI here

Four risks, none of them about the technology failing.

A review with no published test

DWP's defence of the advances model is its human review. "The ultimate safeguard in place is that there is always a human intervention and decision, with no automated decision making by the model." The review is designed with care. The caseworker "does not receive a risk rating in the referral nor are they made aware the referral has been generated by the model", and model referrals are mixed with random control-group referrals "to mitigate against human bias". That is a deliberate guard against the best-documented failure of human oversight, which is deferring to the machine.

It protects the review from the model. It does not show that the review catches what the model gets wrong. The effectiveness figure DWP publishes is a measure of the model. Nothing published measures the caseworker's decision: how often a second review would agree with it, whether it corrects the over-referral of particular groups, or how it fares at appeal.

A fairness test that cannot see

DWP analyses disparity only where at least 70 per cent of claimants have answered the equality questions. Among claimants referred as high risk, that threshold was not met. DWP found no appropriate data source for disability, sex or ethnicity. The protected characteristics it could test were age and nationality. A fairness test that cannot see disability cannot clear the model of disability bias.

Drift between assessments, and a change of name

The 2024-25 assessment found the model "around 3 times more effective" than random; 2025-26 found 2.5 times. The non-UK referral ratio rose from 2.27 to 3.35. The 2024-25 document was titled a *Fairness Assessment*; the 2025-26 document is an *Effectiveness Assessment*, and refers to "last year's Effectiveness Assessment". The renaming is reported here as observed, without inference about why.

Summaries without sources

The Home Office's evaluation of its asylum AI pilots, published in April 2025, found that case summarisation could save 23 minutes per case. It also found that "a small proportion of summaries produced (9%) were deemed to be inaccurate or had missing information" and that "23% of users reported they were not fully confident in the summary information". By design for the pilot, the summaries carried no references back to the transcript. The Home Office describes its findings as reflecting "small-scale pilots" and self-reporting.

WHAT THE TRANSPARENCY REGISTER DOES NOT SHOW

On 29 September 2026 the register held 152 records. Every title was read. The Home Office has one, covering visit visa routing. There is no record for Asylum Case Summarisation or Asylum Policy Search, although the department has told Parliament that policy search is available to all asylum decision-makers. The standard applies to tools within its scope; this briefing states what the register holds, not whether either tool is in scope.

DWP, *Effectiveness Assessment*, 9 July 2026, paras 6, 10, 11, 21, 26 and statistical annex · DWP, *Fairness Assessment*, 17 July 2025 · Home Office, *Evaluation of AI trials in the asylum decision making process*, 29 April 2025 · GOV.UK search, algorithmic transparency records, 29 September 2026.

The threat side

AI is now on both sides of the same decision, and most of what arrives is not fraud.

In insurance, the threat this series found was fraud against the firm. In central government the evidence points somewhere else. The National Audit Office names three forms of external AI-driven activity that public bodies should expect: "an increase in legitimate demand" as the friction of applying falls; "more repeat, low-value or poor-quality submissions"; and "an elevated risk of fraud". The first of those is people reaching services that were previously hard to reach. It is not wrongdoing, and a department that plans for it as a threat will make the wrong decisions about the people it serves.

On the applicant's side

In *UK and R (Munir) v Secretary of State for the Home Department*, promulgated on 17 November 2025, the Upper Tribunal recorded a "considerable increase in the latter half of 2025 in the citation of fictitious authorities" in immigration cases, and held that representatives who put false citations before it should ordinarily expect to be referred to their regulator. That is the professional side of the system using AI badly, and the tribunal has responded.

On the department's side — as reported

In July 2026, The Guardian reported an Upper Tribunal decision on the asylum claim of a Moroccan woman and her child. The Home Office's refusal had cited a country policy note that the judge found appears never to have existed: "It appears that no such CPIN has ever existed." The judge said the letter "bears hallmarks consistent with the use of artificial intelligence", and the First-tier Tribunal had relied on the same document. The Home Office said that claims "are always considered on their individual merits and based on the latest evidence". The allegation is unproven, the case continues, and the decision was not on the tribunal's published database when it was searched for this briefing.

WHAT THE TWO CASES HAVE IN COMMON

Each document went out under a professional's or a department's name. In each, the error was caught later, at the tribunal.

That is the argument of this briefing in miniature. The human was in the loop on paper. What mattered was whether the human's review did anything.

NAO, *Good practice guide for organisations using AI*, May 2026, p.7 · *UK and R (Munir) v SSHD* [2026] UKUT 00081 (IAC) · The Guardian, 28 July 2026; UKAuthority, 30 July 2026.

The UK legal position

The law now turns on one word – meaningful – and on a code the ICO has not yet written.

Section 80 of the Data (Use and Access) Act 2025 replaced Article 22 of the UK GDPR from 5 February 2026. Significant decisions based solely on automated processing moved from being prohibited, with exceptions, to being permitted with safeguards: the person must be told, able to make representations, able to obtain human intervention, and able to contest the decision. Article 22A defines the line: "a decision is based solely on automated processing if there is no meaningful human involvement in the taking of the decision", and in judging that, account must be taken of "the extent to which the decision is reached by means of profiling".

UK GDPR Art. 22A–22D	In force 5 February 2026. Solely automated significant decisions permitted with safeguards. "Meaningful human involvement" decides whether a decision is solely automated.
Art. 22B – special category data	Health and other special category data only with explicit consent, or where the decision is required or authorised by law and the substantial public interest condition applies. That is a condition government can create for itself, by legislating.
Equality Act 2010, s.149	The public sector equality duty. DWP says the advances model has "due regard" to it.
ATRS	Mandatory for central government and in-scope arm's length bodies.
ICO draft ADM guidance	Consulted on from 31 March to 29 May 2026; not yet final.
SI 2026/425	In force 12 May 2026. Requires the Information Commissioner to prepare a code of practice on AI and automated decision-making.

The overlooked change: the code will be evidence

SI 2026/425 is made under section 124A of the Data Protection Act 2018. Section 92 of the Data (Use and Access) Act 2025, in force since 20 August 2025, brought codes prepared under section 124A within section 125, and so within section 127. Section 127 says three things. Failing to follow the code "does not of itself make that person liable to legal proceedings". The code "is admissible in evidence in legal proceedings". And a court, a tribunal or the Commissioner must "take into account a provision of a code" where it is relevant to the question being decided.

Central government is appealed against at a scale no other sector sees: 20,000 benefit appeals were decided at hearing in a single quarter. Not every appeal will turn on how a decision was reached. But where one does, the ICO's code will be the measure of good practice – and the code is not yet written. The ICO has said it will publish timings "in due course". A department that can show its review is meaningful before the code arrives will not be making that case for the first time in front of a judge.

UK GDPR Arts 22A–22D, as substituted by Data (Use and Access) Act 2025 s.80 · DUAA 2025 s.92 (in force 20 August 2025, S.I. 2025/904) · Data Protection Act 2018 ss.124A, 125, 127 · SI 2026/425 · ICO, 23 June 2026 · Equality Act 2010 s.149.

Beyond the UK

Both of this briefing's cases would be high-risk in the EU, and a Dutch court has already been here.

The EU Artificial Intelligence Act lists, in Annex III point 5(a), AI systems "intended to be used by public authorities or on behalf of public authorities to evaluate the eligibility of natural persons for essential public assistance benefits and services", including to "grant, reduce, revoke, or reclaim" them, as high-risk. Point 7(c) adds AI used "to assist competent public authorities for the examination of applications for asylum, visa or residence permits". Under Article 27, public bodies deploying high-risk systems must carry out a fundamental rights impact assessment before use.

The Digital Omnibus on AI, Regulation (EU) 2026/1744, adopted on 8 July 2026 and in force from 27 July, moved the start of the stand-alone high-risk obligations to 2 December 2027.

EU ANNEX III 5(A)

Benefit eligibility, including recovery. An advances fraud model that selects claims for refusal would be assessed against it.

EU ANNEX III 7(C)

Assisting the examination of asylum applications. Case summarisation and policy search would be assessed against it.

The litigation signal

In February 2020 the District Court of The Hague held that the legislation behind SyRI, a Dutch system that combined government data to risk-score people for welfare fraud, breached Article 8 of the European Convention on Human Rights. The Convention is the same one the Human Rights Act 1998 gives effect to in the UK.

A COMPARATOR, NOT AN OBLIGATION

The EU AI Act does not apply to UK departments making decisions about people in the UK. No compliance gap is implied. The value is in seeing where another legislature has drawn the line on the same decisions.

Regulation (EU) 2024/1689, Annex III points 5(a) and 7(c), Article 27 · Regulation (EU) 2026/1744 · Council of the EU, press release, 7 May 2026 · District Court of The Hague, 5 February 2020 (SyRI).

Today

Seven things that can be done now, none of which waits on new law, new money or a supplier.

- **Test the review, not only the model.** Check a sample of caseworker decisions on referred cases against a second review, and publish the agreement rate beside the model's effectiveness figure.
- **Close the equality-data gap before the next assessment.** A fairness test that cannot see disability, sex or ethnicity cannot clear a model of bias on any of them.
- **Publish a transparency record for every tool that shapes a decision,** including tools described as aids to the decision-maker.
- **Put the source beside every AI summary a decision-maker reads,** so the check is one click, not a reread of the transcript.
- **Treat tribunal overturns as evidence about the first decision,** and trace them back to any AI involved in it.
- **Write down the legal basis before any decision touching health data is automated,** under Article 22B.
- **Take the NAO's questions to the audit and risk assurance committee.** The guide was written for them, and asks whether data protection, ethics and legal obligations "are being met in practice".

None of these requires a decision about whether to use more AI. They are the conditions under which the AI already in use can be defended. The first is the one that matters most, because it is the one that turns an assertion about human review into evidence.

Drawn from DWP's assessments of the advances model, the Home Office evaluation of April 2025, the NAO good practice guide (May 2026), PAC reports and MoJ tribunal statistics.

Tomorrow

Three levels defined by where accountability sits, and three horizons.

Three levels

Assistive	It prepares, summarises and surfaces; the decision-maker decides. Case summarisation, policy search and drafting in Copilot sit here. Accountability is unchanged, and almost all defensible value sits here today – provided the summary shows its source.
Augmentive	It selects or recommends within bounds; a person confirms or overrides. The advances model sits here: it chooses which cases a caseworker reviews. The review has to be demonstrably meaningful, and demonstrated rather than asserted.
Adaptive	It decides within a defined envelope, and oversight moves to whoever sets and monitors the envelope. Since February 2026 a solely automated significant decision is lawful with safeguards; for health data, only with consent or a legal basis. For central government this is a reasoned projection, not an observation.

Three horizons

0 to 12 months	The ICO's AI code in preparation. DSIT's delayed commitments to the Public Accounts Committee due by December 2026. DWP's retrained model awaiting testing.
12 to 36 months	The ICO code issued and admissible in evidence. EU high-risk obligations for benefit and asylum systems apply from 2 December 2027.
36 months and beyond	Routine cases decided automatically, with people reviewing exceptions. Whether that review is meaningful is the question this briefing raises.

The third band is a projection, not an observation. The ICO has published no date for its code.

The capacity to get there

In May 2026 DSIT gave the Public Accounts Committee later dates for five of its recommendations on AI in government. Putting senior digital leaders on every executive team moved from June to December 2026. Using government's spending power in AI procurement moved from January to September 2026. Measuring the cost of legacy technology and tracking remediation funding moved from January to autumn 2026. When the committee reported in March 2025, about 28 per cent of central government systems met the definition of legacy, and 21 of the 72 highest-risk legacy systems lacked remediation funding. The NAO names the workforce risks as entry-level learning, "overreliance on automated systems" and "loss of institutional knowledge".

ICO, 23 June 2026 · DSIT, letter to the Public Accounts Committee, 19 May 2026 · PAC, *Use of AI in Government*, 26 March 2025 · NAO good practice guide, May 2026 · Council of the EU, 7 May 2026.

Benefit realisation

Fraud stopped is a real benefit. It is not the whole answer, because it says nothing about the review.

The worked example is on the public record. The NAO estimates the advances model saved £4.4m between May 2022 and its October 2025 report. DWP estimates fraud and error on advances at between £20m and £90m, on £700m paid in 2025-26. Referred advances that are later approved are paid a median of one day later than those approved automatically. Those are the right first measures, and DWP publishes them.

They are measures of the model. The benefit claimed for the model rests on the review, because DWP's case for fairness is that a person always decides. A benefit case that counts fraud prevented will report success whatever the review is doing. The measures below add the review to the case.

WHAT IS ALREADY COUNTED

Money saved. Effectiveness against random. Referral and correct-referral ratios by age and nationality.
Median payment delay.

WHAT WOULD SHOW THE REVIEW WORKS

Agreement between caseworker decisions and a second review. Ratios on a fixed comparator each year.
Appeals and overturns traced to referred cases.
Disparity for disability, sex and ethnicity once the data exists.

The practical step is small. Add one measure of the review itself to every AI business case in the department. It is the measure a tribunal, and in time the ICO's code, will ask about first.

NAO, 22 October 2025 · DWP, *Effectiveness Assessment of Universal Credit Advances Model*, 9 July 2026.

A way in

Eight workstreams. The first three form a natural first step.

01	Readiness assessment	Every AI tool that touches a decision about a person, including the ones described as aids, and who answers for each.	4–6 weeks
02	Use case identification and triage	Sorted by where accountability sits on the three levels, not by expected saving.	3–4 weeks
03	Regulatory and legal position paper	Articles 22A–22D, section 127 and the coming code, the equality duty and the transparency standard, read together.	3–4 weeks
04	Board and executive governance	What the audit and risk assurance committee sees, and how often, including a measure of the review.	4 weeks
05	Threat and integrity response	Rising legitimate demand, AI-drafted submissions and AI-drafted decisions, handled as three different things.	4 weeks
06	Benefit realisation framework	Including the review measures in section eleven.	3 weeks, parallel
07	Build, buy or partner evaluation	Including where the error lands when a supplier's tool is wrong.	4–6 weeks
08	Twelve-month mobilisation roadmap	Sequenced against the ICO's code and the department's own assessment cycle.	2 weeks

THE FIRST STEP

Workstreams one to three form a natural first engagement of six to eight weeks.

The ICO is preparing its AI code now. Six to eight weeks is enough to know whether a department's review would stand up to it — which is the difference between describing a practice you already have and being told what it should have been.

nc@neilcatton.com · neilcatton.com

Sources

Every figure in this briefing, traceable to a named publication and a date.

- Department for Work and Pensions, *Effectiveness Assessment of Universal Credit Advances Model*, 1 April 2025 to 31 March 2026, published 9 July 2026.
- Department for Work and Pensions, *Universal Credit Advances Model – Fairness Assessment*, 1 April 2024 to 31 March 2025, published 17 July 2025.
- National Audit Office, *Good practice guide for organisations using AI*, May 2026.
- National Audit Office, *Tackling benefit overpayments due to fraud and error*, 22 October 2025.
- Committee of Public Accounts, *Use of AI in Government*, Eighteenth Report of Session 2024–25, 26 March 2025.
- Committee of Public Accounts, press release on its report on fraud and error, 27 March 2026.
- Department for Science, Innovation and Technology, letter from the Permanent Secretary to the Chair of the Public Accounts Committee, 19 May 2026.
- Ministry of Justice, *Tribunal Statistics Quarterly: April to June 2026*, 10 September 2026.
- Home Office, *How many people are in the UK asylum system?*, Immigration system statistics, year ending June 2026, 27 August 2026.
- Home Office, *Evaluation of AI trials in the asylum decision making process*, 29 April 2025.
- House of Commons, written answer UIN 121674 (Home Office), 24 March 2026.
- GOV.UK, algorithmic transparency records: all 152 titles, retrieved 29 September 2026.
- *UK and R (on the application of Munir) v Secretary of State for the Home Department* [2026] UKUT 00081 (IAC), promulgated 17 November 2025.
- The Guardian, 28 July 2026, and UKAuthority, 30 July 2026, reporting an Upper Tribunal decision of 14 July 2026 (decision not located on the published database).
- UK GDPR Articles 22A–22D, as substituted by the Data (Use and Access) Act 2025, section 80, in force 5 February 2026.
- Data (Use and Access) Act 2025, section 92; Data Protection Act 2018, sections 124A, 125 and 127; SI 2026/425.
- Information Commissioner's Office, consultation on draft guidance on automated decision-making, 31 March–29 May 2026; statement of 23 June 2026.
- Parliamentary and Health Service Ombudsman, *Our strategy 2026 to 2031*.
- Equality and Human Rights Commission, *An update on our approach to regulating artificial intelligence*, 30 April 2024.
- Equality Act 2010, section 149.
- Regulation (EU) 2024/1689, Annex III and Article 27; Regulation (EU) 2026/1744; Council of the EU press release, 7 May 2026.
- District Court of The Hague, judgment of 5 February 2020 (SyRI).

Excluded on the record: government's £45bn "full potential digitisation" estimate and its £36bn automation component; the 62% "routine tasks" figure; the NAO's 2023 adoption survey, as dated; unsourced figures on DWP bank monitoring; a campaign-commissioned legal opinion on asylum AI; a secondary report of case summarisation's full rollout; DWP's reported-illness ratios, where the comparator is named inconsistently; and any year-on-year age comparison.

About the author

Neil Catton is an independent technology strategist with more than thirty-nine years of technology leadership across twenty-plus sectors — from central government, policing and national health systems to banking, manufacturing and insurance. He has held CTO and Director-level roles at BT, Wipro, Fujitsu, Sopra Steria and Capita, and works across a fractional portfolio spanning public safety, public sector, insurance, health technology and early-stage startups.

In government technology he was Chief Technology Officer for Home Affairs and Criminal Justice at Fujitsu, serving the Home Office, the Ministry of Justice, HMCTS and borders and immigration. At Sopra Steria he was Sector CTO for government. At Capita he led the UK application architecture practice — more than a hundred architects working across central and local government, justice, health and defence. He is the author of the technology ethics trilogy *The Next Evolution*, *The Cognitive Crucible* and *The Shadow System*, and writes at neilcatton.substack.com.

How this briefing was produced

Only primary sources are used: a department on its own published record, an auditor, a select committee, a tribunal, official statistics or the statute book. Where this briefing reports an absence, the relevant source was searched for it and the claim is stated as what was searched. Figures that could not be traced to a named source are named as excluded rather than quietly dropped, because what a briefing leaves out is the better test of it.

In keeping with this report's own argument: this briefing was produced with AI assistance, used transparently and kept under direct human editorial control throughout. Every source was verified; every judgement was reviewed by Neil Catton before publication.

NOTICE

This briefing is not legal, regulatory, financial or professional advice. The position is stated as at 29 September 2026 and may since have moved; readers should check the current position before relying on anything here. Third-party sources are cited in good faith and should be verified before reliance. Named organisations appear only from their own publications, audit and parliamentary records, or tribunal findings, and are used illustratively; none has been engaged or approached in connection with this briefing. No advisory or client relationship is created by receipt of this document.

Independence. The author holds no vendor affiliations, reseller agreements, referral arrangements, commission interests or equity positions in anything referenced in this briefing, and has received no payment from any organisation named in it. That independence is what makes the exclusion discipline described above credible rather than merely claimed.

© Neil Catton 2026. Internal use and circulation permitted with attribution.

Version history

Version	Change	Date
v1.0	First publication.	29 September 2026